

Writing an Incomplete Shell

2014-15, CSCI 3150 - Programming Assignment 1

Specification version 1.0, 2014 Sep 11

Updated: 23:00, 2014 Oct 2

Abstract

The purpose of this assignment is to introduce students to the existing useful system calls in the Linux environment and the control of the creation as well as the termination of the processes.

Contents

1	Introduction	3
2	Program Flow	4
2.1	Input command line	5
2.2	Command line interpreter	6
2.3	Process creation	8
2.3.1	Locating the programs	8
2.3.2	Wildcard expansion	9
2.3.3	I/O Redirection	10
2.4	Detecting process termination and suspension	12
2.5	Built-in shell commands	12
2.6	Signal Handling	14
2.7	Incomplete job control	15
2.7.1	Definition of job and job control	15
2.7.2	Suspension handling	15

2.7.3	List the suspended jobs	16
2.7.4	Resuming the suspended jobs	16
2.7.5	Defects in the incomplete job control	17
2.8	Extra requirements	18
3	Milestones and Deliverables	19
3.1	Phase 1: command line interpreter (6%)	19
3.1.1	Input and Output	19
3.1.2	Requirements for Phase 1	21
3.1.3	Deliverables	22
3.2	Phase 2: the complete shell (12%)	22
3.2.1	Overview	22
3.2.2	Phase 2 milestones	22
3.2.3	Deliverables	23
3.3	Marking	23

1 Introduction

A significant portion of the interactive actions of Unix/Linux system is performed through an user-level command interpreter known as the **shell**, e.g., *bash* and *tcsh*. The shell implements many important aspects of the Linux operating system in terms of your daily use, including *process management* and *I/O redirection*.

In this assignment, you are required to implement a primitive version of the shell which makes heavy uses of the system calls provided by Linux.

What are the things that you are supposed to learn from this assignment?

- Hard skills:
 - Understand the internal workings of a shell;
 - Write a medium-sized program (500 - 1,000 lines of codes);
 - Design and implement a structured program (and hope that you still remember what a structured program is);
 - Write and compile a program with more than one source file (if you want to);
 - Read man pages and dig out every little detail before you use a system/function call;
 - Using useful, existing function/system calls;
- Soft skills:
 - Manage your time under a *tight schedule*;
 - Cooperate with your partner *in harmony*;
 - Teach your partner with *patience*;
 - Learn from your partner *with an open mind*;

2 Program Flow

The shell program that you are going to implement is a simpler version than the ones that you use in Unix or Linux system. For simplicity, we name the shell that you are going to implement as the *OS shell*. Figure 1 shows the execution flow of the OS shell.

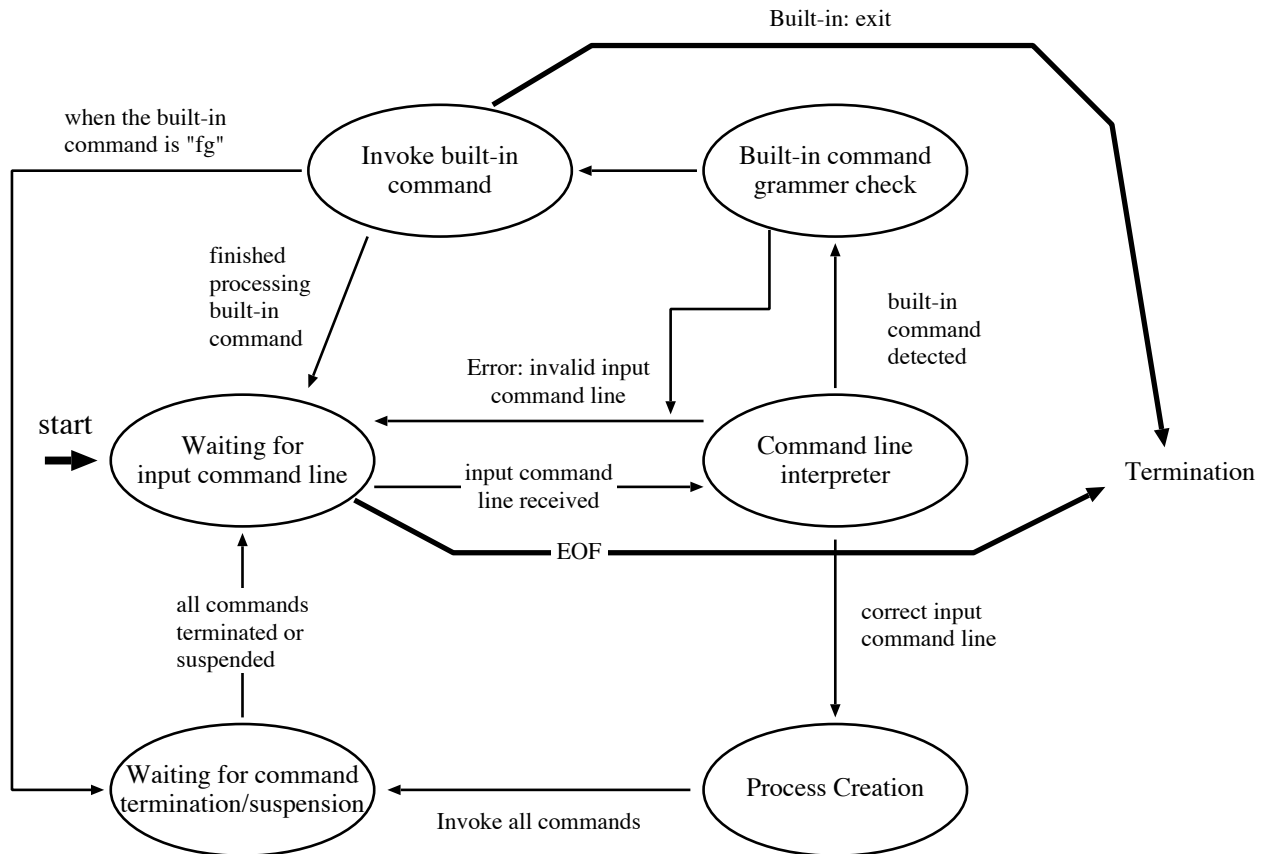


Figure 1: The shell's execution flow.

1. When the OS shell first starts, the program should be in the state “**Waiting for input command line**”. The OS shell should show the following:

```
[3150 shell:/home/tywong]$ _
```

We usually call it **the prompt**. Note that “/home/tywong” is an example directory at which the shell is invoked.

2. When the user types in a command followed by a **carriage return** (or the “enter” key), e.g.,

```
[3150 shell:/home/tywong]$ /bin/ls | /bin/more
```

Then, the received input will be processed by the “**Command line interpreter**” of the OS shell, and we called the input, “/bin/ls | /bin/more”, the **input command line**.

- If the command line interpreter finds that the input command line **agrees with the pre-defined syntax**, which will be defined in Section 2.2, then the OS shell should invoke the commands specified by the input command line.
- Else, the input command line will not be executed and the OS shell waits for another input command line, i.e., going back to the state “*Waiting for input command line*”.

2.1 Input command line

The **input command line** is a character string. The OS shell should read in the input command line from the **standard input stream** (`stdin` in C) of the OS shell. To ease your implementation, there are assumptions imposed on the input command line:

Assumptions
<ol style="list-style-type: none">1. An input command line has a maximum length of 255 characters.2. An input command line ends with a carriage return character ‘<code>\n</code>’.3. There will no leading or trailing space characters in the input command line.4. A token is a series of characters without any space characters. Each token (or word) in the input command line is separated by at least one space character.

2.2 Command line interpreter

The command line interpreter checks whether the input command line matches a pre-defined language or not.

- If yes, the OS shell checks and invokes the program(s) specified in the input command line.
- Else, the OS shell waits for the next input command line.

Before we know whether a given input command line is correct or not, we need to know the language that the input command line should follow. The language that the shell takes is similar to the *bash* shell and is specified as follows:

```
[start]  := [empty string] or
          := [built-in command] [arg]* or
          := [command] [recursive] or
          := [command] [terminate]
```

```
[recursive] := '|' [command] [recursive] or
             := '|' [command] [terminate]
```

```
[terminate] := [empty string]
```

```
[command]  := [command name] [args]*
[command name] := A string without any space, tab, > (ASCII 62), < (ASCII 60), | (ASCII 124), * (ASCII 42), ! (ASCII 33), ' (ASCII 96), ' (ASCII 39), nor " (ASCII 34) characters. Note that the command name is assumed to be different from a built-in command.
[arg]      := A string without any space, tab, > (ASCII 62), < (ASCII 60), | (ASCII 124), ! (ASCII 33), ' (ASCII 96), ' (ASCII 39), nor " (ASCII 34) characters.
```

Ease-your-pain restrictions.

Maximum number of pipes: we restrict the maximum number of pipes, i.e., “ ”, in a command line to two only.

Some examples of valid command line inputs are as follows.

- blank line
- Simple command: “cat”, “cat *”, and “cat dog”
- Pipeline: “cat | cat”, and “cat | cat | cat”

Some examples of invalid command line inputs are as follows.

- “cat | cat | ”
- “>_<” (using invalid characters.)
- “cat | cat | cat | cat” (due to the pipe restriction.)

Important: you are not required to ensure that a command works perfectly. You only need to check the input command line against the grammar only.
--

Every valid input command line should begin with the [start] rule. Any input command line that does not agree with the above language is considered to be wrong. The following shows are the requirements of the command line interpreter:

1. If an input command line violates the above language, the shell has to return an message “**Error: invalid input command line**” before prompting for the next command.
2. If an input command line agrees with the above language, the OS shell will enter the “**Process creation**” state.

Useful Functions.

fgets(), strtok()

2.3 Process creation

When the input command line is valid, the OS shell should **create the processes** according to the names specified in the input command line.

2.3.1 Locating the programs

There are two kinds of paths:

- If a path's name begins with the character '/', then that path is an **absolute path**.
- Else, that path is a **relative path**.
- Names starting with "./" and "../" are considered as absolute paths though they are not started with the character '/'.

If a program name is specified by an **absolute path**, e.g., /bin/ls or ./a.out, then the creation of the process is straight-forward. Else, if a command name is specified by a **relative path**, then the OS shell has to search for the program from the following locations in the following order:

Required path-searching order.
/bin → /usr/bin → . (current directory)

Remember that the **first matched entry** stops the search.

Example.
There exists two files: "/bin/ls" and "./ls". Then, the command "ls" will cause the OS shell to invoke the program "/bin/ls".

If the specified program cannot be found in the above three locations, the OS shell should report "[command name]: command not found". We repeat the previously illustrated example as follows.

Example.
[3150 shell:/home/tywong]\$ /bin/sl /bin/more /bin/sl: command not found [3150 shell:/home/tywong]\$ _

Note that:

- The `exec` system call family is able to report whether you have failed to invoke a command or not by looking at its **return value** and **error number**, `errno`.

In this assignment, you only need to report the cases that a program cannot be found (with both absolute and relative paths). If the `errno` specifies other errors, then the shell should report:

`[command name]: unknown error`

- **Hint:** You are not required to perform an actual search for the command's location. You can totally depend on the `execve()` call family.
- Searching for commands outside the above three paths is prohibited. You will risk a **mark deduction** if your OS shell can invoke programs that are outside the above three paths.

Useful resource.
<code>fork()</code> system call, <code>exec</code> system call family, <code>setenv()</code>

2.3.2 Wildcard expansion

The wildcard in our shell is denoted by the character `'*'` only. In this assignment:

- A wildcard character will only appear in an argument to a program.
- In a command line, there can be more than one argument carrying a wildcard.
- An argument can contain more than one wildcard character.
- A wildcard can appear in the following ways:
 - `"/bin/ls *.txt"`: in the beginning of a token;
 - `"/bin/ls hell*txt"`: in the middle of a token;
 - `"/bin/ls hello*"`: in the end of a token.

The meaning of wildcard expansion is that you are replacing the wildcard expression by looking up the files that matches that wildcard expression.

Example.	
Say there are four files in the directory: “a.txt”, “b.txt”, “c.txt”, and “abc.mp3”	
<code>ls *</code>	\Rightarrow <code>ls a.txt abc.mp3 b.txt c.txt</code>
<code>ls *.txt</code>	\Rightarrow <code>ls a.txt b.txt c.txt</code>
<code>ls a*</code>	\Rightarrow <code>ls a.txt abc.mp3</code>
<code>ls b* a*</code>	\Rightarrow <code>ls b.txt a.txt abc.mp3</code>
<code>ls *.iso</code>	\Rightarrow <code>ls *.iso</code>

The expansion of wildcard expressions is actually **a task of the shell**. Therefore, in your shell implementation, when you find a token with wildcard characters, you have to:

1. Treat the token as one wildcard expression.
2. Then, expand that expression into a list of tokens, and this list of tokens should be sorted lexicographically.
3. When you cannot expand a particular wildcard expression, just **keep the expression unexpanded** (as shown in the last example).

Note importantly that you do not need to implement the expansion by yourself. Please use the library call “`glob()`”. Our tutor will cover this in the tutorials.

2.3.3 I/O Redirection

You are only required to implement one **I/O redirection feature** in this assignment: the pipe. We usually name a command line with pipes a pipeline and the following is a basic rule in representing a pipeline.

Pipe, basic rule
<code>[command 1] [command 2]</code>

A pipe connects the standard output stream of “[command 1]” to the standard input stream of “[command 2]”. In other words, the standard input stream of “[command 2]” is replaced by the standard output stream of “[command 1]”.

Pipe, example
<pre>[3150 shell:/home/tywong]\$ cat output.txt first hello world second hello world [3150 shell:/home/tywong]\$ cat output.txt head -1 first hello world [3150 shell:/home/tywong]\$ _</pre>

where “head -[n]” works similar to cat, but it prints only the first n lines from its standard input stream to its standard output stream, unlike cat which prints all data.

Useful system calls for I/O redirection
pipe().

Please note the following important points about the pipeline:

- First thing first, every process within a pipeline must be executing concurrently.
- When one of the commands in the pipeline fails to execute, the shell should:
 1. report the corresponding error message(s), and then
 2. terminate that corresponding process only.

On the other hand, the shell should let the remaining processes execute normally. As a consequence, the remaining processes should terminate gracefully if and only if the processes are connected through the pipes correctly.

2.4 Detecting process termination and suspension

Once the processes are launched, the shell has to wait until **all the processes have stopped executing**, either terminated or suspended. After the processes have stopped running,¹ the OS shell should wait for the next input command line.

Useful system calls.
<code>waitpid()</code> only. <code>wait()</code> cannot help.

Note important that the OS shell should not leave any **zombies** in the system when the OS shell is ready to read a new input command line. Otherwise, **you will have 1% of the course mark deducted**.

2.5 Built-in shell commands

In a traditional shell, there are built-in shell commands. In this assignment, we implement **four** of them, and their requirements are listed as follows.

1. **cd** [**arg**]. This command is called “*change directory*” and it changes the current working directory of the shell. Some points to note:
 - The argument [**arg**] is the destination path to which the user wants to change. Note that [**arg**] can be either an absolute path or a relative path.
 - If the location specified by [**arg**] results in an error, the OS shell should report the error message: “[**arg**]: cannot change directory”.
 - In addition, this built-in shell command takes only one argument. If the number of arguments is wrong, the OS shell should report the error message: “**cd: wrong number of arguments**”.

Useful system calls. <code>chdir()</code> .

¹A process can stop running because of normal or abnormal reasons.

2. **exit**. This command terminates the shell. This command takes no arguments. If the number of arguments is wrong, the OS should report: “**exit: wrong number of arguments**”. Note that if there are suspended jobs, the OS shell will not be terminated.

“exit” example, with suspended jobs

```
[3150 shell:/home/tywong]$ exit - are you sure?
exit: wrong number of arguments
[3150 shell:/home/tywong]$ cat
^Z
[3150 shell:/home/tywong]$ exit
There is at least one suspended job
[3150 shell:/home/tywong]$ _
```

3. **fg [job number]**. This command is called “*foreground*”, and it wakes a suspended job, specified by the job number - **[job number]**. We will talk about this built-in command in detail in Section 2.7 on Page 15.

Note that, this built-in shell command takes only one argument. If the number of arguments is wrong, the OS shell should report the error message: “**fg: wrong number of arguments**”.

4. **jobs**. This command prints a list of suspended jobs to the standard output stream of the OS shell.

“jobs” example

```
[3150 shell:/home/tywong]$ jobs
[1] cat
[2] cat | cat
[3] top
[3150 shell:/home/tywong]$ _
```

The jobs are sorted by the time that they are created for the first time. The smallest job number is always one, and the OS shell should not skip any numbers. We will talk about this built-in command in detail in Section 2.7.

Important

The priority of the built-in shell commands always comes first. In other words, suppose there exists a program `/bin/cd`, then when the user types in `cd`, only the built-in shell command `cd` will be invoked instead of `/bin/cd`.

2.6 Signal Handling

Though a shell plays a special role in Unix/Linux operating systems, as a matter of fact, it is just a normal user-space program. Therefore, the OS shell will handle every possible signal in its default manner. However, you will not expect that a shell will be stopped or terminated by signals such as **Ctrl + Z** and **Ctrl + C**. Therefore, you are required to change the signal handling routine of the OS shell. The signals that you should take care of are listed as follows:

Signal	Signal Handling
SIGINT (Ctrl + C)	Ignore the signal.
SIGTERM (default signal of command “kill”)	Ignore the signal.
SIGQUIT (Ctrl + \)	Ignore the signal.
SIGTSTP (Ctrl + Z)	Ignore the signal.
SIGSTOP	Default signal handling routine.
SIGKILL	Default signal handling routine.

For signals that are not listed in the above table, the OS shell should adopt the default signal handling routines of the corresponding signals. Note importantly that only the OS shell changes the handling of the said signals. That means the child processes created by the OS shell should, however, adopt the default signal handling routines.

Signal handling example

```
[3150 shell:/home/tywong]$ cat
^C
[3150 shell:/home/tywong]$ _
```

In the above example, the signal **SIGINT** is generated by **Ctrl + C**, and that signal terminates the process “**cat**” only, but not the OS shell.

2.7 Incomplete job control

You are going to implement an subset of job control features, comparing to those of the shell programs you can find in Unix/Linux.

2.7.1 Definition of job and job control

What is a job? A job is created by a **valid input command line**. All the processes that are created by the OS shell, because of a valid input command line, are considered as one job. Job control means the way that we control all those processes together. In the OS shell, we implement the suspension and the continuation of the jobs.

2.7.2 Suspension handling

When a job is running, e.g., “**cat**”, the user can suspend it through the signal “**SIGTSTP**”, which can be generated by typing **Ctrl + Z**. In brief, there are three ways to suspend a process.

1. Type “**Ctrl + Z**” on the terminal in which the process is running;
2. Send the signal “**SIGTSTP**” by using */bin/kill*. E.g., */bin/kill -TSTP [process name]*;
3. Use the system call `kill()`. Yet, you have to write a program to invoke the system call “`int kill(pid_t pid, int signal_number)`”, with the signal **SIGTSTP**.

When a job is suspended, the OS shell should go back to the “**Waiting for command termination/suspension**” state (see Figure 1 on page 4). The following example illustrates this point.

Job suspension, example
[3150 shell:/home/tywong]\$ cat cat ^Z [3150 shell:/home/tywong]\$ _

2.7.3 List the suspended jobs

When jobs created by the OS shell become suspended, the OS shell should be able to list those jobs using the **built-in shell command** “jobs”.

List suspended job, example

```
[3150 shell:/home/tywong]$ jobs
No suspended jobs
[3150 shell:/home/tywong]$ cat | cat
^Z
[3150 shell:/home/tywong]$ jobs
[1] cat | cat
[3150 shell:/home/tywong]$ _
```

2.7.4 Resuming the suspended jobs

The OS shell should also provide another built-in shell command “fg” for the user to resume the suspended jobs.

Wake suspended job, example

```
[3150 shell:/home/tywong]$ jobs
[1] cat | cat
[3150 shell:/home/tywong]$ fg 1
Job wake up: cat | cat
^C
[3150 shell:/home/tywong]$ jobs
No suspended jobs
[3150 shell:/home/tywong]$ _
```

Note that when the supplied job number does not exist, you should report the error message: “fg: no such job”. Note also important that the OS shell is not required to handle the case that a suspended job is **killed** by other processes.

Useful resource
signal() system call (read the man-page using the command “man 7 signal” on Linux and “man -s 3HEAD signal” on Unix) and waitpid() system call.

2.7.5 Defects in the incomplete job control

A frustrating defect would appear when you are playing with signals in a *wrong* way.

- [The wrong signal]. A strange scenario happens when “Ctrl + C” is pressed while there are suspended jobs.
- [The wrong thing(s) appeared]. The suspended jobs are killed!

Defect example:
[3150 shell:/home/tywong]\$ cat ^Z [3150 shell:/home/tywong]\$ cat ^C [3150 shell:/home/tywong]\$ fg 1 Job wake up: cat [3150 shell:/home/tywong]\$ _ //woooo...my poor cat...

Hence, you do not have to worry about the above strange behavior because this is the defect of the design of this assignment. (Of course, writing more codes can solve the problem. But, the lecturer would like to stop torturing you.)

If you are questing for perfection, please look up the following keywords from Google and Wikipedia and have a perfect job control implementation:

- Process group, foreground process group, background process group.
- Library call: `tcsetpgrp()`, system call: `setpgid()`, `getpgid()`.

2.8 Extra requirements

The following requirements are restrictions over the executions of programs. The **rule of thumb** of the following restrictions is to ensure that your shell is implemented

- through your own command line interpreter, and
- by using the fork-execute-wait(pid) system call combination.

In other words, you should experience how to build a shell using system calls. Most importantly, the restrictions ensure that credits are only given to those who have spent efforts.

1. **[Extra requirement #1]** You are not allowed to invoke the `system(3)` library call. Otherwise, you would score 0 marks for this assignment.
2. **[Extra requirement #2]** You are not allowed to invoke any existing shell programs in this assignment. Otherwise, you would score 0 marks for this assignment. One exception is that the input command lines is to invoke those existing shell programs.

Note that the phrase “existing shell programs” implies the shell programs installed in the operating system including, but not restricted to, `/bin/sh` `/bin/bash`, etc.

3 Milestones and Deliverables

This assignment contributes 18% of the total marks of the course, and we break this assignment into **2 phases**. You must use either C or C++ to implement your assignment. **Otherwise, you would receive zero marks.**

3.1 Phase 1: command line interpreter (6%)

You have to implement a correct command line interpreter. Referring back to Figure 1, you have to implement from “Start” to “Termination” through the following nodes and edges:

- Node: “Waiting for input command line”.
- Node: “Command line interpreter”.
- Node: “Built-in command grammar check”.
- Partial implementation of Node: “Invoke built-in command”. You have to complete the implementation of the following built-in commands: `cd` and `exit`;

For `jobs` and `fg`, you are required to have their corresponding grammar check only. In Phase 2, you will complete the implementation of their functionalities.

- Edge: “input command line received”.
- Edge: “built-in command detected”.
- Edge: “finished processing built-in command”.
- Edge: “Error: invalid input command line”.
- Edge: “EOF”.

3.1.1 Input and Output

1. **Input.** The interpreter reads input command lines from the standard input stream (`stdin` in C). The program should stop when it encounters the end of the standard input, i.e., the “Ctrl + D” input, or the `exit` built-in command.

2. **Output - for correct inputs.** If the syntax of the input command line is correct, the interpreter should break the input into **tokens** and print out the type of the tokens line by line. The output format is given as follows:

Phase 1 output format	
Token 1:	"Content" (Type)
Token 2:	"Content" (Type)
...	
Token n:	"Content" (Type)

where “Content” is the token itself and “Type” can be either one of the following:

Phase 1 output format, “Type”			
Built-in Command	Argument	Pipe	Command Name

Phase 1 input and output example
<pre>[linux3:/home/tywong]\$./phase1 [3150 shell:/home/tywong]\$ ls Token 1: "ls" (Command Name) [3150 shell:/home/tywong]\$ cat * Token 1: "cat" (Command Name) Token 2: "*" (Argument) [3150 shell:/home/tywong]\$ cat >_< Error: invalid input command line [3150 shell:/home/tywong]\$ cd / / cd: wrong number of arguments [3150 shell:/home/tywong]\$ cd / Token 1: "cd" (Built-in Command) Token 2: "/" (Argument) [3150 shell:/]\$ exit Token 1: "exit" (Built-in Command) [Shell Terminated] [linux3:/home/tywong]\$ _</pre>

Note that in Phase 1, you are not required to implement a complete shell. Let us consider the examples shown above:

- [No process execution]. “ls” command line will not be executed although the command line passed the grammar check.
- [Detecting built-in commands’ grammar]. “cd / /” yields an error because the number of arguments “cd” is restricted to be one.
- [No wildcard expansion]. We will leave this feature to Phase 2.
- [Execute correct “cd” and “exit”]. The examples:
 - cd /
 - exit

are correct command lines and the OS shell must execute them.

3. **Output - for incorrect inputs.** If the syntax of the input command line is incorrect, the interpreter should print the message “**Error: invalid input command line**” to the standard output stream (`stdout` in C).

3.1.2 Requirements for Phase 1

- The interpreter recognizes built-in commands whenever a correctly-placed built-in command is encountered.
- The interpreter is not required to create processes, to execute the programs, nor to recognize if the programs could be located.
- The interpreter should terminate when either the end of the standard input stream or the `exit` built-in command is encountered.
- To facilitate a fast grading process, in Phase 1, you have to print the results by **accurately following the format stated in the above example output**.

3.1.3 Deliverables

You are required to submit a set of source files. **Never submit any executables.** The files should be compiled into one executable only.

Deadline: 23:59, Sep 28, 2014 (Sunday).

3.2 Phase 2: the complete shell (12%)

With a working interpreter, the shell can make use of the result from the interpreter to create processes. We break the programs into milestones (marks are attached, too). Notice that those milestones are **incremental**. E.g., Milestone 2 is building on top of Milestone 1. You are recommended to finish the milestones in the order they are presented.

3.2.1 Overview

Task	Marks	M1	M2	M3	M4
Commands with arguments	2%	•	•	•	•
With pipes	5%		•	•	•
Signals	1%			•	•
Job control + Remaining built-in commands	4%				•

* Without loss of generality, “M1” means *Milestone 1*.

3.2.2 Phase 2 milestones

- **[Milestone 1]**. This requires you to implement a simple shell which can create processes and execute commands with arguments. You should integrate the command line interpreter from Phase 1 into the shell.
- **[Milestone 2]**. This requires you to implement the support for the pipes. Remember, there are **at most two pipes** in a command line.
- **[Milestone 3]**. This requires you to handle signals as instructed in Section 2.6 on Page 14.

- [Milestone 4]. Last but not least, you have to implement the job control features described in Section 2.7 on Page 15 together with all built-in commands.
- Note that the wildcard expansion inputs will be tested over all the above milestones.

Since the milestones are incremental by design, marks would be lost if a bug from an earlier milestone appears when you are showing us a latter milestone.

Say you had passed the tests for Milestone 2, i.e., pipes. However, a hidden bug in handling pipes manifested while you were demonstrating the job control features. Then, you would lose marks for the concerned testcase in the job control milestone, but not the marks for the pipe features.

3.2.3 Deliverables

You are required to submit a set of source files. **Never submit any executables.** The files should be compiled into one executable only. Therefore, you should better plan ahead so that your phase 1 program can smoothly integrated into phase 2.

~~Deadline: 23:59, Oct 12, 2014 (Sunday).~~

Deadline: 23:59, Oct 26, 2014 (Sunday).

3.3 Marking

The marking is separated based on the two phases mentioned.

- The marking of the interpreter program will be separated from the shell program.
- There can be cases that the program you submitted in Phase 1 is incorrect. Then, you may ask: *“how can I finish the assignment with a buggy interpreter?”*

Don’t worry. We allow the update of the interpreter’s codes. Please just put the new code together with the shell program in the same submission package. As a matter of fact, we will mark your assignment as follows.

- We will grade your interpreter from your phase 1 submission.
 - Then, we grade your shell from your phase 2 submission.
 - We will not re-grade your interpreter from phase 2.
 - If bugs appeared during the grading of the shell, we would never distinguish the source of the bugs (from phase 1 or not). We consider them as incorrect implementations of the involved features of the shell.
- We grade your assignment based on the input and the output. We seldom look into your code, locate your faults, nor give you partial marks.

Last but not least, we will go through days of **demonstrations** in grading your submissions. You and your partner both have to attend the demonstration and instruct our tutors how to work with your submissions. This is to make sure that you can appeal for the results right at the spot. Please stay tune with the course homepage about the dates of the demonstration.

Submission Guideline

For the submission of the assignment, please refer to the following link:

<http://www.cse.cuhk.edu.hk/csci3150/>

Last reminder:

Phase 1 Deadline: 23:59, Sep 28, 2014 (Sunday).

~~**Phase 2 Deadline: 23:59, Oct 12, 2014 (Sunday).**~~

Phase 2 Deadline: 23:59, Oct 26, 2014 (Sunday).

–END–

Change Log

Time	Details
23:00, 2014 Oct 2	Phase 2 deadline is changed to “23:59, Oct 26, 2014 (Sunday)”.